

Introducing the Quantum Research Kernels: Lessons from Classical Parallel Computing

A. Y. Matsuura

Applications and Quantum Architecture Lab

Intel Corp.

Hillsboro, OR, USA

anne.y.matsuura@intel.com

Timothy G. Mattson

Parallel Computing Lab

Intel Corp.

Ilwaco, WA, USA

timothy.g.mattson@intel.com

Abstract—Quantum computing represents a paradigm shift for computation requiring an entirely new computer architecture. However, there is much that can be learned from traditional classical computer engineering. In this paper, we describe the Parallel Research Kernels (PRK), a tool that was very useful for designing classical parallel computing systems. The PRK are simple kernels written to expose bottlenecks that limit classical parallel computing performance. We hypothesize that an analogous tool for quantum computing, Quantum Research Kernels (QRK), may similarly aid the codesign of software and hardware for quantum computing systems, and we give an example of a representative QRK.

Index Terms—Quantum Computing, HW/SW co-design, Parallel computing, Architecture

I. INTRODUCTION

Quantum computers have the potential to transform computing. Their advent promises to open up new classes of applications and enable the solution of problems that are intractable for classical computers today. To reach that potential, however, requires great technological innovations, including research advances in fundamental physics, the development of new quantum programming models that “normal” humans can use, and solutions to a host of issues in both software and hardware concerning building quantum systems.

Quantum computing has a long way to go, but we believe we can leverage best practices from classical digital computing to advance through the stages of development at a faster rate. In this paper, we focus on one lesson learned from the era of digital computing: that the best systems emerge from a hardware/software co-design process. All too often hardware is designed and then “thrown over the fence” to software developers to figure out how to make the hardware useful. It is more effective to have the software and hardware teams working together, so the programming systems are ready as soon as the hardware is available and the hardware includes features specifically needed to make the software both more efficient to run and easier to write (for example [1]).

However, it is challenging to do hardware/software codesign when you don’t know what future applications will look like. Hypothesis: Even though we may not know what future applications will look like, we do know the features of a system that will limit these applications. If we can define these “bottlenecks” and build a system that collectively minimizes

their impact, we can be confident that the system will be effective for these future applications.

This was the basic principle behind Parallel Research Kernels (PRK) [2] in the parallel computing field. The PRK are a collection of simple kernels that expose the features of a system that limits parallel performance. They are small, generate their own data, do a computation (so systems can’t “cheat”), and test their results. In essence, they are a way for application programmers to precisely define what they need from a system; to guide hardware developers to build systems that will work well for applications. They have proven useful for designing systems [1] and to explore the suitability of extreme scale programming models [3].

The goal of this paper is to launch a conversation about using an approach similar to the PRK for quantum computing. We call these the Quantum Research Kernels (QRK). Can we define features of quantum computing systems that will limit applications for these systems? Can we produce well defined kernels to expose these features? Can we anticipate the breadth of application design patterns so we can be confident that the set of QRK are complete? These are challenging questions. In this paper, we define the problem and propose a process to answer these questions.

II. PARALLEL RESEARCH KERNELS

The Parallel Research Kernels stress a system in ways parallel applications in high performance computing would. They are listed in Table I where we provide the name of each kernel, a brief definition, and the features of a parallel system stressed by the kernel. The kernels are defined mathematically and with a reference implementation using C, OpenMP, and MPI. They are available in a github repository [4].

We submit that the Parallel Research Kernels are complete. If a system is built that does well with all of them, then it is likely that system will be effective for running parallel HPC applications. The PRK were selected by an ad hoc committee of parallel application programmers. When we started the project (in 2005) parallel computing in various forms had been around over 25 years. We were able to convene a committee of experts with decades of experience. Over the course of several meetings and long email chains, the committee came up with the list of kernels.

TABLE I: **Parallel Research Kernels:** – A set of basic kernels designed to stress features of a system that limit the performance of HPC applications.

Name	Definition	Exposed system feature
Transpose	Transpose a dense matrix	Bisection Bandwidth
Reduce	Elementwise sum of multiple private vectors	Message passing, local memory bandwidth
Sparse	Sparse-matrix vector product	scatter/gather operations
Random	Random update to a table	Bandwidth to memory with random updates
Synch_global	Global synchronization	Collective synch.
Synch_p2p	point-to-point synchronization	message passing latency, remote atomics
Stencil	Stencil method	nearest neighbor and asynchronous comm.
Refcount	Update shared or private counters	Mutual exclusion locks
Nstream	daxpy over large vectors	Peak memory bandwidth
DGEMM	Dense Matrix Product	Peak floating point perf.
Branch	Inner loop with branches	Misses to the instruction cache, branch prediction.
PIC	Particle in cell	unstructured asynch. multitasking
AMR	Adaptive mesh refinement	hierarchical asynch. multitasking

III. QUANTUM RESEARCH KERNELS

The quantum research kernels (QRK), inspired by the PRK, will define a set of kernels designed to expose features of a quantum computing system that will constrain the success of applications written for quantum computers. The QRK are intended to be a sufficiently complete set such that if a system were constructed that did well for each of the QRK, that system would most likely be a successful system for supporting key applications.

As with the PRK, the QRK will be produced through a community-driven process by programmers interested in writing applications for quantum computers. To help nurture this conversation, we provide an example of a potential QRK. One of the bottlenecks for scalable quantum computing, is the difficulty of loading the data (particularly classical data) into the quantum machine. State preparation can be an exponentially hard problem itself, leading to the conundrum that loading of the data into the quantum machine can completely negate the speed-up gained by doing the problem on a quantum computer [5].

Hence, the first QRK follows.

- **Name:** Encode
- **Definition:** Create a sequence of classical values from $i = 0$ to $i = N$ equal to $4i\pi/N$
- **Action:** Encode qubits with the values from the previous step. Rotate each qubit by $\pi/6$.
- **Test:** Read qubits and confirm that they have the correct rotated value

Note this has all the features we would expect in a QRK. It defines problem input that is generated so the QRK can scale to any number of qubits. A specific operation is defined so at the end of the QRK, the result can be validated. Finally, it exposes a specific feature of a quantum computer that will limit the ability of applications to run on the system expressed

in terms architects can understand and use to guide the design of a quantum computer.

We have the beginnings of two additional QRK. The first is one we call the *Computational Area*. This is the product of a number of qubits that can be entangled times the number of operations that can be carried out before the entangled state can no longer be maintained. A high Computational Area can be produced by a small number of qubits that remain entangled over a large number of operations or by having a large number of qubits that remain entangled for a small number of operations. There are advantages to both cases so this measure supports both.

the second QRK we are working on is called *Parallel Streams*. This measures the ability of a quantum computer to execute multiple independent streams of operations at the same time and in parallel. There are a number of options on how to define the work that must be carried out in parallel. To start with, we'd run the Computational Area QRK in each stream, though as we consider a wider range of applications for quantum computers, we may find a better case for each stream.

These three QRK are just a start. We need a larger set that covers the full range of features needed from a successful quantum computer, hopefully, on the order of 10. The system features each QRK stresses overlap, but taken together they need to cover the full range of features need by application programmers from a quantum computer so those designing parallel systems can be confident a system that does well on the full set of QRK will meet the needs of applications programmers. This is the primary goal of the QRK. However, once established, an application designer can model an the needs of an application in terms of a linear combination of QRK thereby using them to help select the quantum computer best suited to a particular application.

IV. CONCLUSION

We believe that the quantum computing field can learn from classical computing design techniques. The PRK are a powerful design tool, and we posit that there is a need for a similar approach for hardware/software co-design for quantum computing. We have introduced the concept of QRK and have provided a few examples of QRK. This paper is a call to the quantum computing user community to work together to develop a complete set of QRK that can guide the design and development of quantum computing.

REFERENCES

- [1] T. Mattson, R. van der Wijngaart, M. Riepen, T. Lehnig, P. Brett, W. Haas, P. Kennedy, J. Howard, S. Vangal, N. Borkar, G. Ruhl, and S. Dighe, "The 48-core scc processor: the programmer's view," in *ACM/IEEE conference on Supercomputing*, 2010.
- [2] R. F. Van der Wijngaart and T. G. Mattson, "The parallel research kernels," in *High Performance Extreme Computing Conference (HPEC), 2014 IEEE*. IEEE, 2014, pp. 1–6.
- [3] R. F. V. der Wijngaart, A. Kaya, J. Hammond, G. Jost, T. S. John, S. Sridharan, T. Mattson, J. Abercrombie, and J. Nelson, "Comparing runtime systems with exascale ambitions using the parallel research kernels," *Proceedings International Supercomputing Conference*, 2016.
- [4] "Parallel Research Kernels (PRK)," <https://github.com/ParRes/Kernels>.

- [5] S. Aronson, “Read the fine print,” *Nature Physics*, vol. 11, pp. 291–293, 2015.